

# [CAPSTONE] Fruits Freshness Classifier based on Comprehensive Image Features

## Executive Summary

### Problem Statement:

- It is difficult to identify the freshness of the fruits one by one when the number of fruits is in large volumes
- It is important to monitor the freshness of the fruits from time to time to ensure good customer satisfaction and retention

### Scope:

- The dataset covers fruits specifically on apples, bananas, and oranges

### Definitions:

- Total RGB Color Composition: The cumulative measurement of color distribution across the red, green, and blue channels within each image.
  - Red: Total Colour on the Red Channel
  - Green: Total colour on the Green Channel
  - Blue: Total colour on the Blue Channel
- Picture Intensity (Contrast): A component of texture analysis, indicating the degree of brightness or darkness in the image.
- Overall Energy: A measure of the aggregate visual intensity and vibrancy present in the image.
- Pixel Correlation: An assessment of the interdependence and relationship between adjacent pixels in the image.
- Homogeneity: An evaluation of the uniformity and consistency of visual elements throughout the image.

### Methodology:

- Use Accuracy, Precision, Recall (Sensitivity), F1, MCC and Kappa to evaluate the classification performance of the models
- Use RMSE, MSE, MAE to evaluate the regression performance of the models

## Importing the Libraries

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from imblearn.over_sampling import SMOTE

# Scaler Methods
from sklearn.preprocessing import RobustScaler, MaxAbsScaler, MinMaxScaler

# Machine Learning Models

from sklearn.model_selection import cross_val_score, train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier

# Metrics
from sklearn.metrics import make_scorer
from sklearn.metrics import (
    make_scorer,
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    matthews_corrcoef,
    cohen_kappa_score,
    mean_squared_error,
    mean_absolute_error,
    confusion_matrix,
    ConfusionMatrixDisplay,
    classification_report
)

import pickle
```

## Importing the Database

```
In [ ]: df = pd.read_csv("fruits.csv")
```

```
In [ ]: df
```

Out[ ]:

	Image	Red	Green	Blue	Contrast	Energy	Cc
0	rotated_by_75_Screen Shot 2018-06-08 at 5.25.0...	480016	320784	308173	1455.458367	0.091725	
1	rotated_by_15_Screen Shot 2018-06-08 at 5.15.2...	429908	287132	294462	1898.266122	0.089732	
2	rotated_by_15_Screen Shot 2018-06-08 at 5.26.4...	493084	503330	363093	1304.415102	0.298530	
3	rotated_by_15_Screen Shot 2018-06-08 at 5.27.0...	398133	463717	277134	1057.761224	0.087325	
4	rotated_by_75_Screen Shot 2018-06-08 at 5.33.4...	403308	424642	264396	2111.405306	0.109698	
...	...	...	...	...	...	...	...
10896	saltandpepper_Screen Shot 2018-06-12 at 11.46....	510687	494520	443423	252.043265	0.143389	
10897	rotated_by_75_Screen Shot 2018-06-12 at 11.40....	337028	264447	225432	882.749388	0.131601	
10898	Screen Shot 2018-06- 12 at 11.46.17 PM	514385	463767	426552	441.551020	0.113603	
10899	rotated_by_15_Screen Shot 2018-06-12 at 11.45....	452977	368142	305902	1738.992653	0.224891	
10900	saltandpepper_Screen Shot 2018-06-12 at 11.23....	604274	481870	322724	222.860816	0.200812	

10901 rows × 9 columns

# Exploratory Data Analysis (EDA)

In [ ]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10901 entries, 0 to 10900
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Image           10901 non-null  object
1   Red             10901 non-null  int64
2   Green           10901 non-null  int64
3   Blue            10901 non-null  int64
4   Contrast        10901 non-null  float64
5   Energy          10901 non-null  float64
6   Correlation     10901 non-null  float64
7   Homogeneity     10901 non-null  float64
8   Freshness       10901 non-null  int64
dtypes: float64(4), int64(4), object(1)
memory usage: 766.6+ KB
```

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	Red	Green	Blue	Contrast	Energy
<b>count</b>	10901.000000	10901.000000	10901.000000	10901.000000	10901.000000
<b>mean</b>	466876.253646	402949.903036	337048.713788	1337.803641	0.25819
<b>std</b>	83431.591490	85979.015033	84021.923106	777.397036	0.14067
<b>min</b>	74212.000000	54894.000000	38685.000000	31.596735	0.01622
<b>25%</b>	418183.000000	345000.000000	279333.000000	647.005714	0.15970
<b>50%</b>	472741.000000	400620.000000	333381.000000	1387.058776	0.23309
<b>75%</b>	522343.000000	460991.000000	392566.000000	1789.657143	0.33182
<b>max</b>	633913.000000	622438.000000	571283.000000	5078.617143	0.78822

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: Image           0
Red               0
Green            0
Blue             0
Contrast         0
Energy           0
Correlation      0
Homogeneity      0
Freshness        0
dtype: int64
```

```
In [ ]: df.dtypes
```

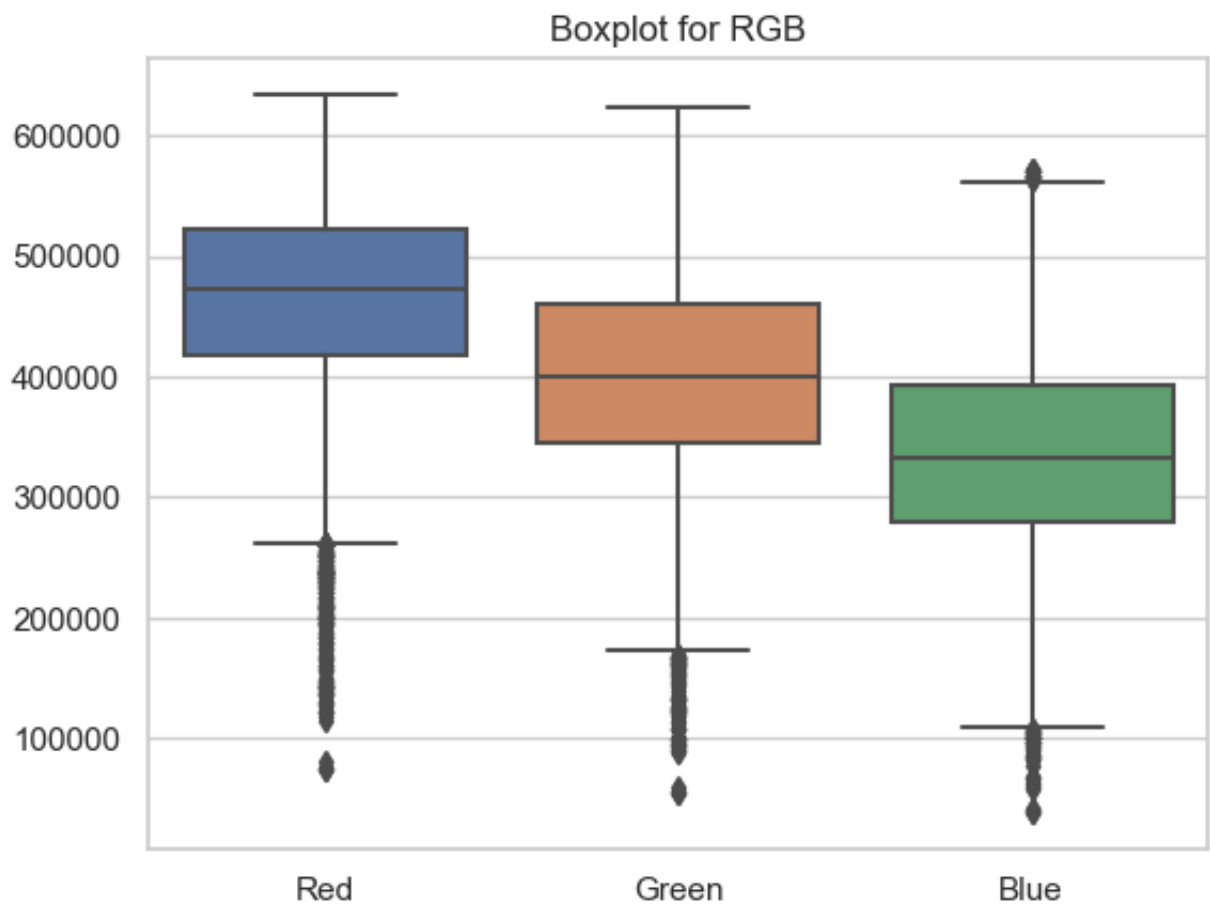
```
Out[ ]: Image          object
        Red            int64
        Green          int64
        Blue           int64
        Contrast       float64
        Energy         float64
        Correlation    float64
        Homogeneity    float64
        Freshness      int64
        dtype: object
```

```
In [ ]: df = df.drop(columns=['Image'])
```

```
In [ ]: def remove_outliers(df, column):
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        return df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]
```

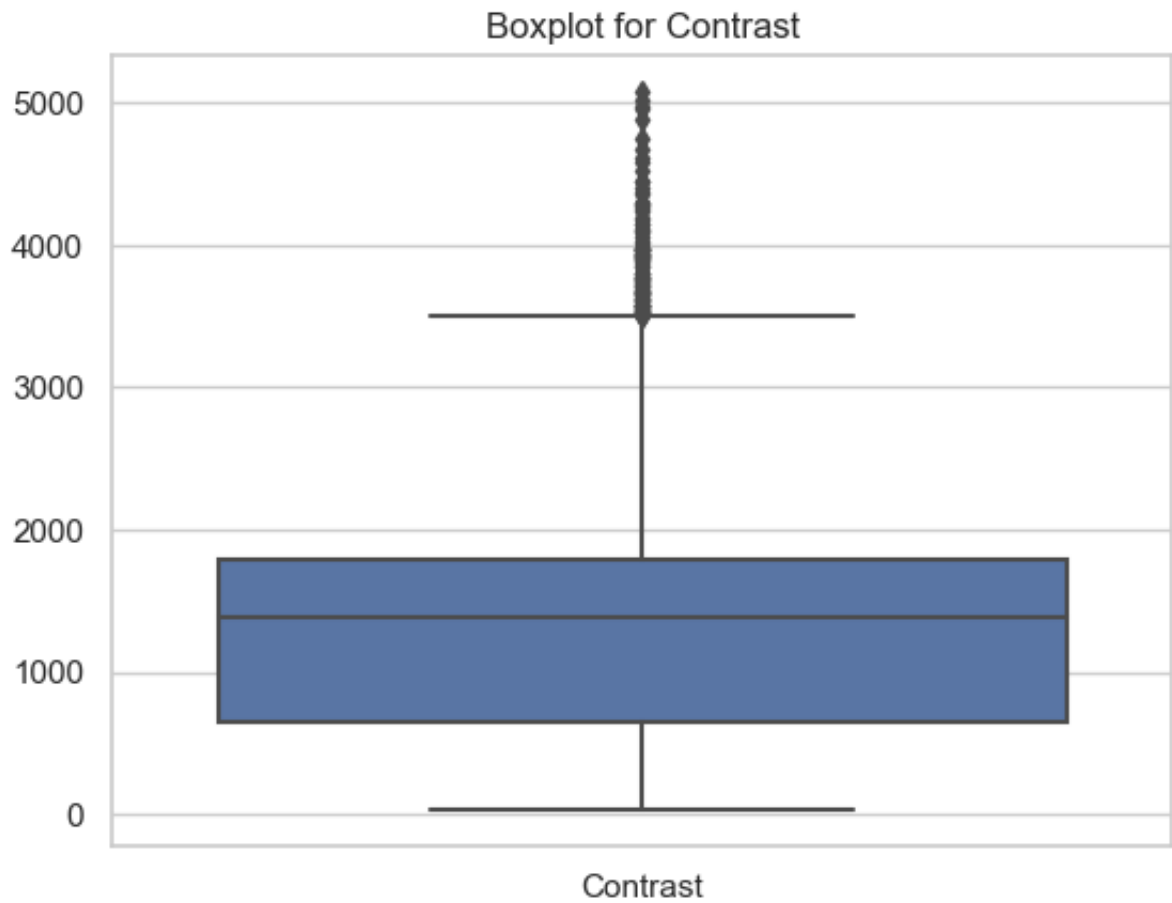
```
In [ ]: columns = ['Red', 'Green', 'Blue']
        df_rgb = df[columns]

        sns.boxplot(data=df_rgb)
        plt.title("Boxplot for RGB")
        plt.show()
```

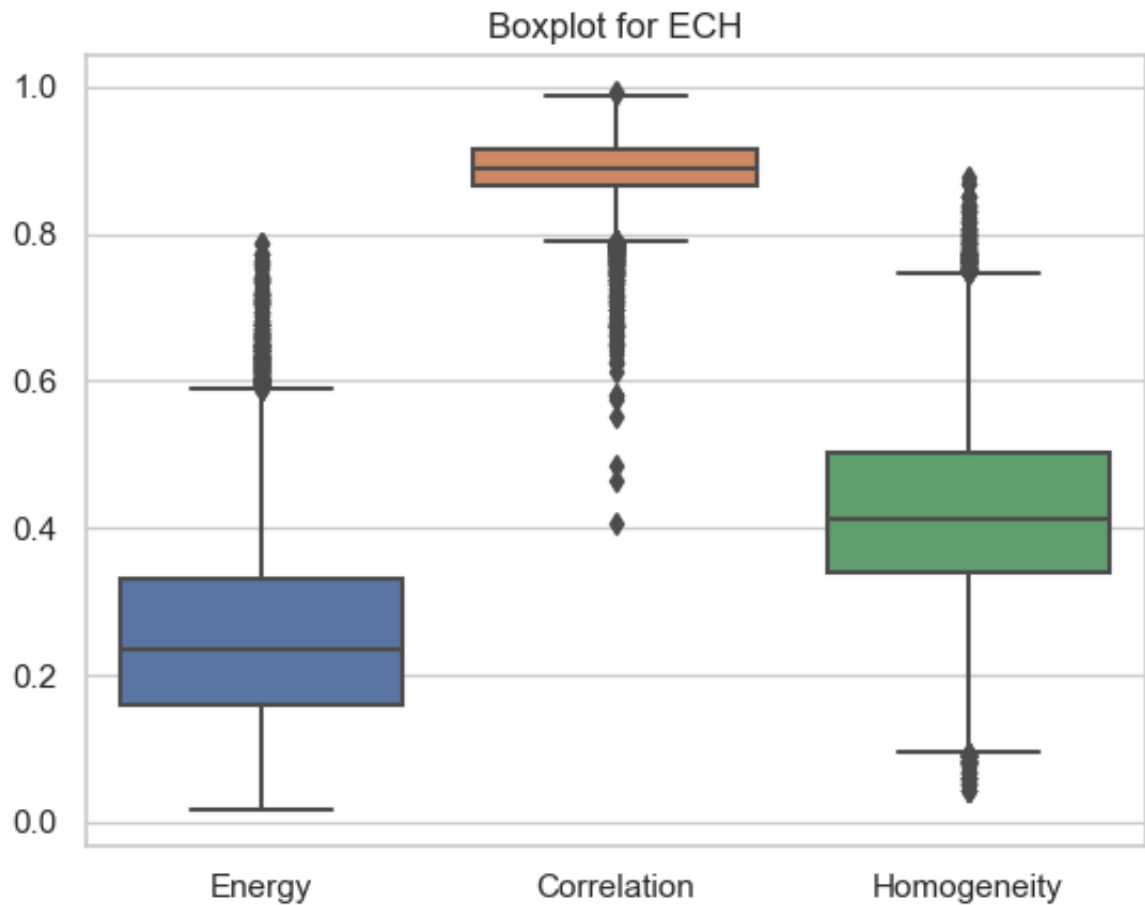


```
In [ ]: columns = ['Contrast']
df_contrast = df[columns]

sns.boxplot(data=df_contrast)
plt.title("Boxplot for Contrast")
plt.show()
```



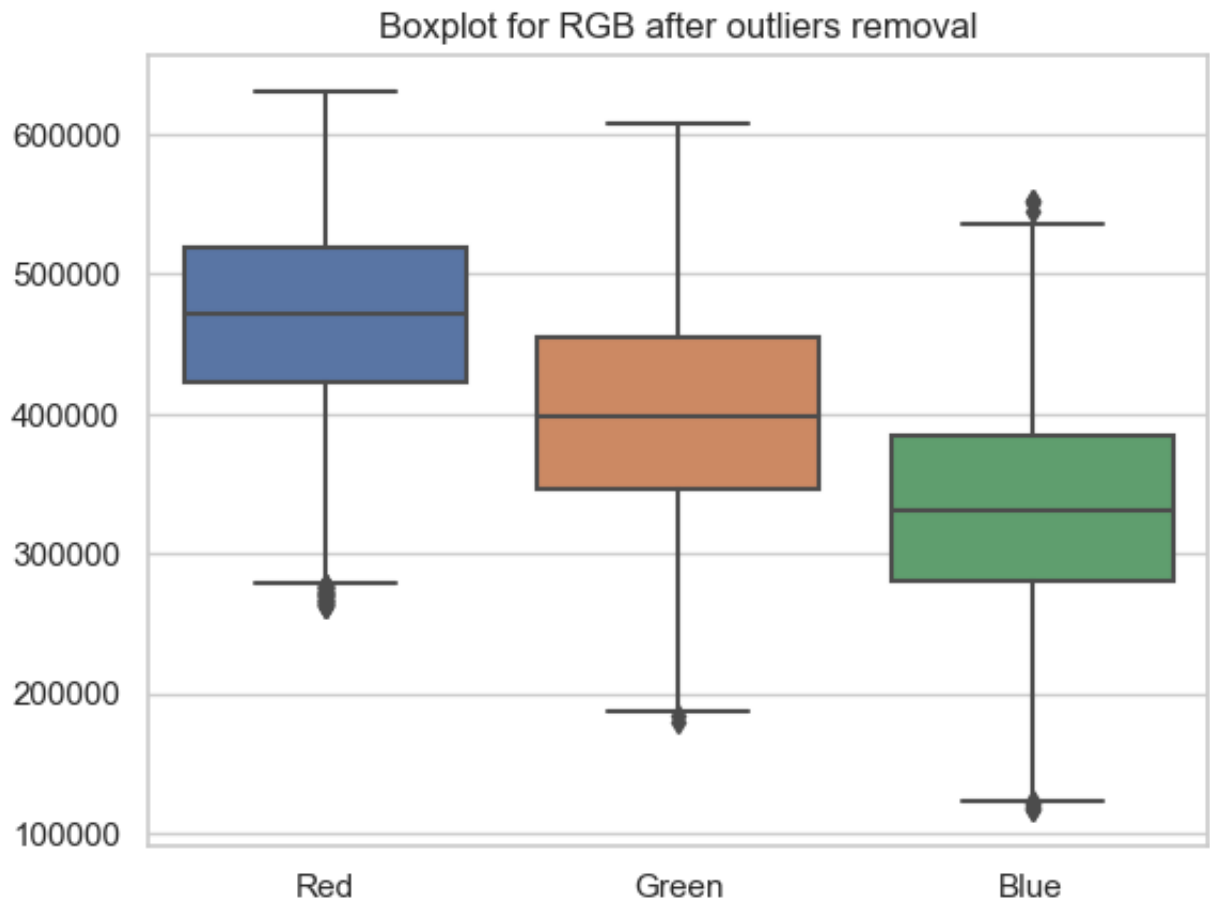
```
In [ ]: columns = ['Energy', 'Correlation', 'Homogeneity']  
df_ech = df[columns]  
  
sns.boxplot(data=df_ech)  
plt.title("Boxplot for ECH")  
plt.show()
```



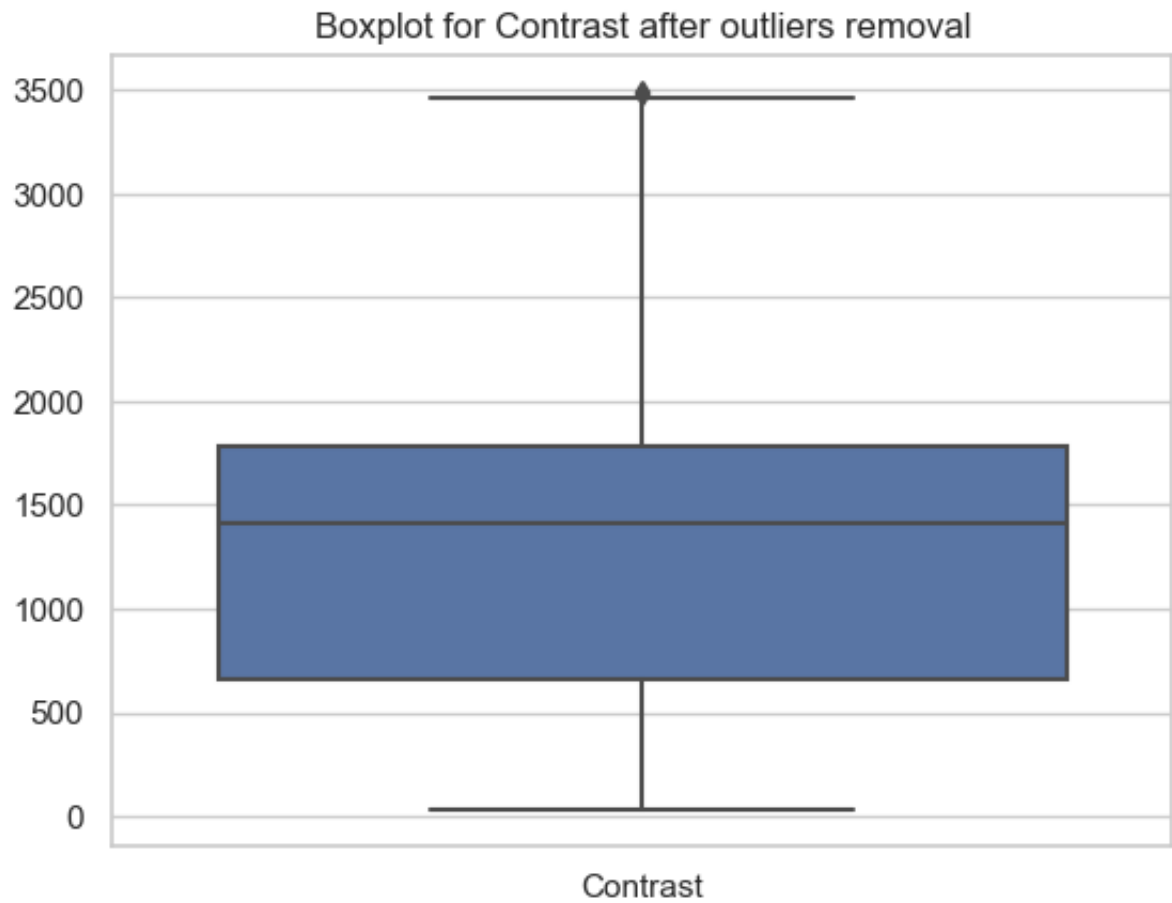
```
In [ ]: for col in df.columns:  
        df = remove_outliers(df, col)
```

```
In [ ]: columns = ['Red', 'Green', 'Blue']  
df_rgb = df[columns]  
  
sns.boxplot(data=df_rgb)  
plt.title("Boxplot for RGB after outliers removal")  
plt.show()
```






```
In [ ]: columns = ['Contrast']  
df_contrast = df[columns]  
  
sns.boxplot(data=df_contrast)  
plt.title("Boxplot for Contrast after outliers removal")  
plt.show()
```



```
In [ ]: columns = ['Energy', 'Correlation', 'Homogeneity']  
df_ech = df[columns]  
  
sns.boxplot(data=df_ech)  
plt.title("Boxplot for ECH after outliers removal")  
plt.show()
```




```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 9764 entries, 0 to 10900  
Data columns (total 8 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   Red              9764 non-null   int64  
1   Green            9764 non-null   int64  
2   Blue             9764 non-null   int64  
3   Contrast         9764 non-null   float64  
4   Energy           9764 non-null   float64  
5   Correlation      9764 non-null   float64  
6   Homogeneity      9764 non-null   float64  
7   Freshness        9764 non-null   int64  
dtypes: float64(4), int64(4)  
memory usage: 686.5 KB
```

## Correlation Matrix

```
In [ ]: corr_matrix = df.corr()  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', square=True)  
plt.figure(figsize=(10, 8))  
plt.show()
```



<Figure size 1000x800 with 0 Axes>

\*\* Since the correlation between the features and freshness are not relatively strong, considering to use non-linear machine learning models here.

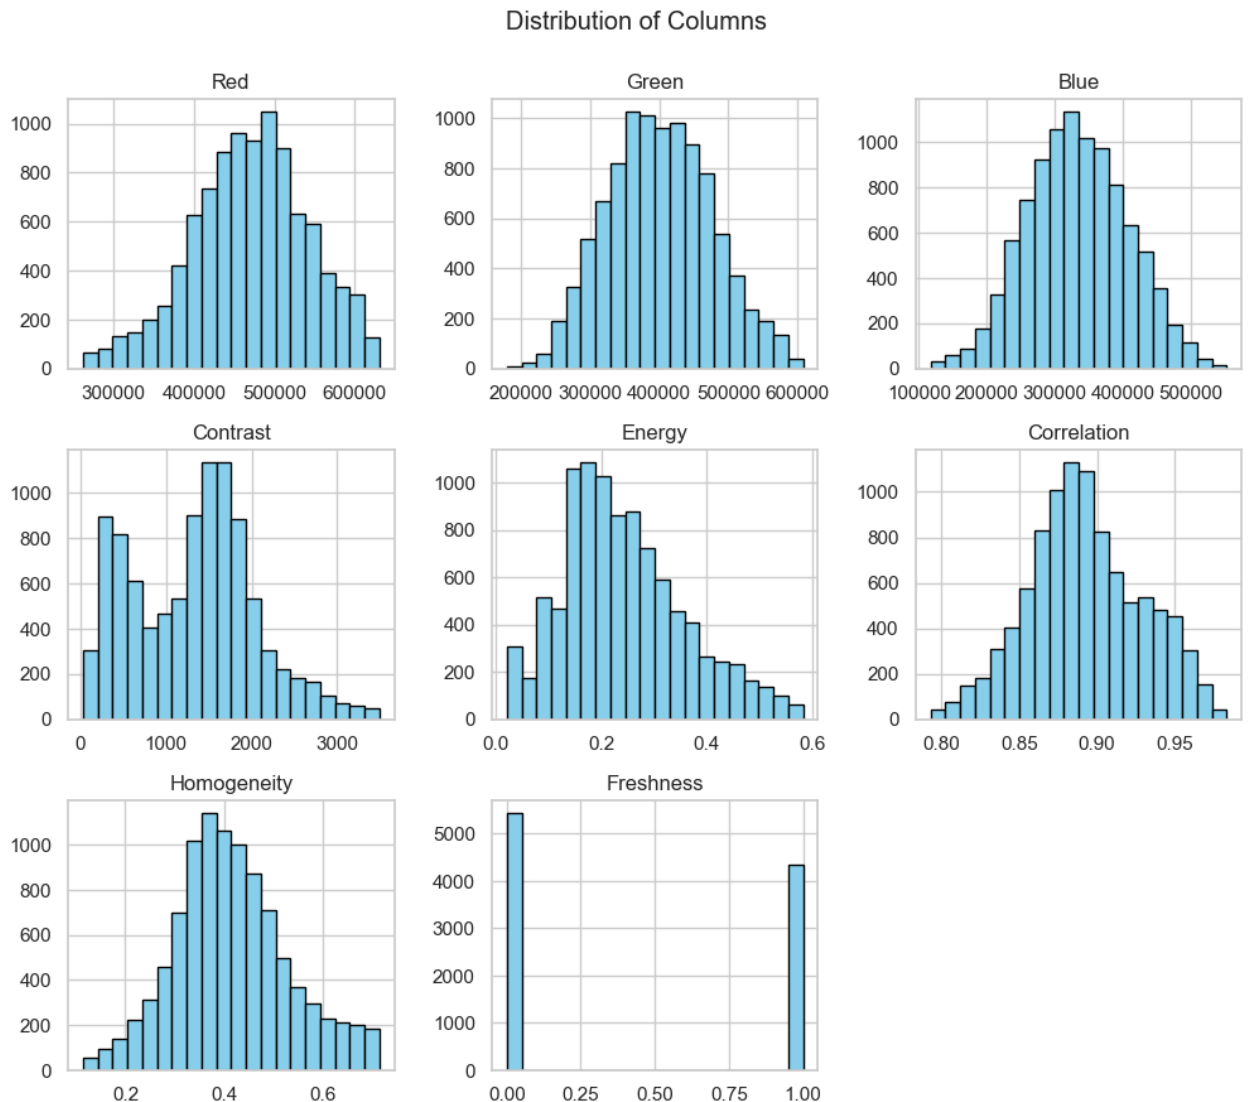
## Distribution for Each Columns

```
In [ ]: skewness = df.skew()

print("Skewness for each column:")
print(skewness)
```

```
Skewness for each column:
Red          -0.219872
Green         0.120736
Blue          0.051449
Contrast      0.238257
Energy        0.553582
Correlation   0.059359
Homogeneity   0.300997
Freshness     0.227631
dtype: float64
```

```
In [ ]: # Assuming df is your DataFrame
df.hist(figsize=(12, 10), bins=20, color='skyblue', edgecolor='black')
plt.suptitle('Distribution of Columns', y=0.95)
plt.show()
```



**\*\* Only Energy is relatively left-skewed**

```
In [ ]: # Filter the DataFrame for fresh and rotten fruits
fresh_df = df[df["Freshness"] == 1]
rotten_df = df[df["Freshness"] == 0]

# List of feature names to visualize
feature_names = ["Red", "Green", "Blue", "Contrast", "Energy", "Correlation"]

# Set up Seaborn styles
sns.set_theme(style="whitegrid")

# Create a plot for each feature with fresh vs. rotten comparison
fig, axs = plt.subplots(len(feature_names), 1, figsize=(10, len(feature_names)))

for i, feature in enumerate(feature_names):
```

```

sns.histplot(fresh_df[feature], kde=True, ax=axes[i], color="green", label=fresh)
sns.histplot(rotten_df[feature], kde=True, ax=axes[i], color="red", label=rotten)
axes[i].set_title(f"Distribution of {feature} for Fresh vs. Rotten Fruits")
axes[i].set_xlabel(feature)
axes[i].set_ylabel("Frequency")
axes[i].legend()

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

perating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```



```
with pd.option_context('mode.use_inf_as_na', True):
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```


```
with pd.option_context('mode.use_inf_as_na', True):
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```


```
with pd.option_context('mode.use_inf_as_na', True):
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```









```
In [ ]: # This two features doesn't differentiate fresh or rotten
feature_names.remove("Energy")
feature_names.remove("Correlation")
```

## Feature Extraction

```
In [ ]: x = df.drop(columns=['Energy', 'Correlation', 'Freshness'])
y = df['Freshness']
```

## Scaling the Data

```
In [ ]: scalers = {
    'Robust Scaler': RobustScaler(),
    'Max Abs Scaler': MaxAbsScaler(),
    'Min-Max Scaler': MinMaxScaler(),
    'Standard Scaler': StandardScaler()
}

x_robust_scale = scalers['Robust Scaler'].fit_transform(x)
x_max_abs_scale = scalers['Max Abs Scaler'].fit_transform(x)
x_min_max_scale = scalers['Min-Max Scaler'].fit_transform(x)
x_standard_scale = scalers['Standard Scaler'].fit_transform(x)

fig, axs = plt.subplots(5, 1, figsize=(8, 15))
axs[0].boxplot(x.values, vert=False, patch_artist=True, boxprops=dict(facecolor='lightblue'))
axs[0].set_title('Original Data')

scaling_results = {
    'Robust Scaler': x_robust_scale,
    'Max Abs Scaler': x_max_abs_scale,
    'Min-Max Scaler': x_min_max_scale,
```

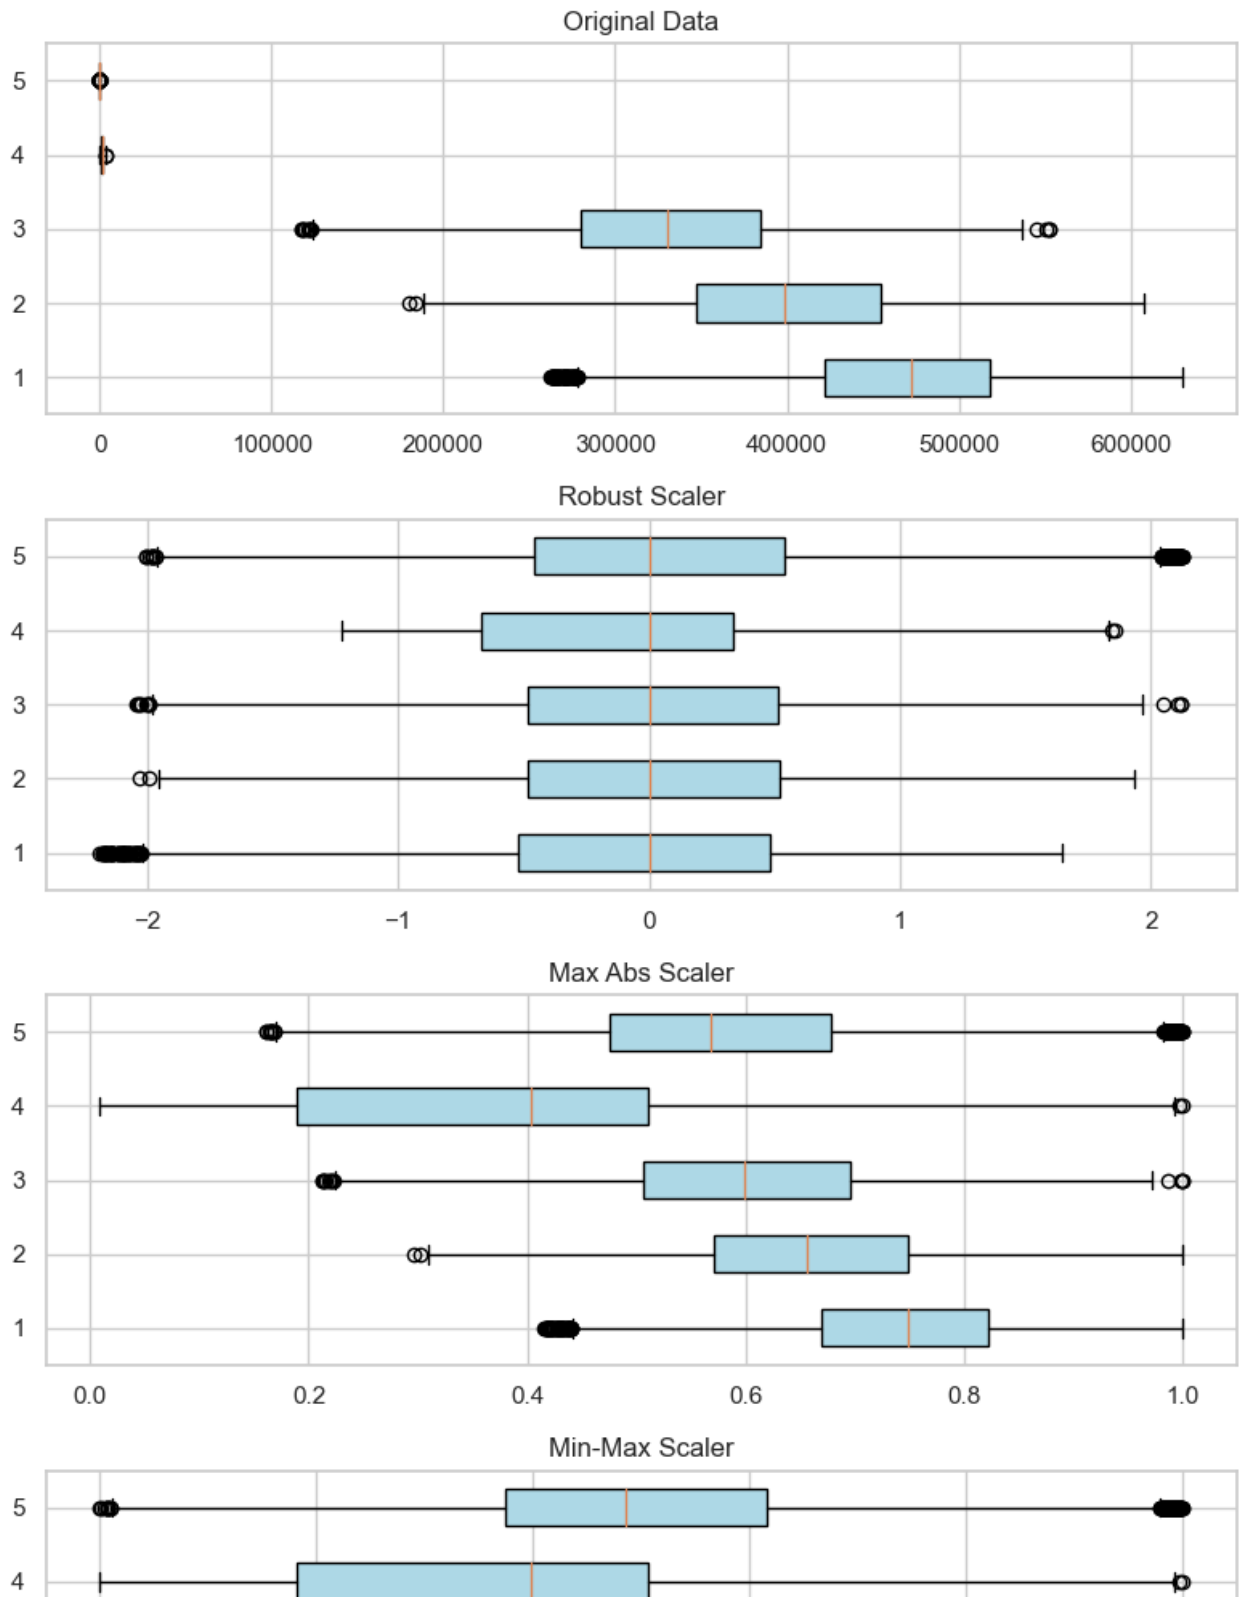
```


'Standard Scaler': x_standard_scale
}

for idx, (scaler_name, scaled_data) in enumerate(scaling_results.items(),
    axes[idx].boxplot(scaled_data, vert=False, patch_artist=True, boxprops
    axes[idx].set_title(scaler_name)

plt.tight_layout()
plt.show()

```





\*\* Although there are significant difference for each of the scaler methods, it is important to know which scaler should be paired with which type of machine learning models \*\* Max Abs Scaler is suitable when the data has positive and negative values, since the data here are all in positive, this option may not be required \*\* Min-max Scaler is suitable for Neural Network Models

## Model Selection & Training

### Linear

- 1. Logistic Regression

### Non-linear

- 2. Support Vector Machine (SVC)
- 3. Decision Tree
- 4. Random Forest
- 5. K Neighbors (KNN)
- 6. Naive Bayes
- 7. XGBoost
- 8. AdaBoost

### Neural Network

```
In [ ]: x_train, x_test, y_train, y_test = train_test_split(x_standard_scale, y,
```

```

num_rotten = (y_train == 0).sum()
num_fresh = (y_train == 1).sum()

print("Number of instances labeled as 'rotten':", num_rotten)
print("Number of instances labeled as 'fresh':", num_fresh)

```

Number of instances labeled as 'rotten': 3812

Number of instances labeled as 'fresh': 3022

```

In [ ]: smote = SMOTE(random_state=42)
x_train, y_train = smote.fit_resample(x_train, y_train)

num_rotten_resampled = (y_train == 0).sum()
num_fresh_resampled = (y_train == 1).sum()

print("Number of instances labeled as 'rotten' after resampled:", num_rotten_resampled)
print("Number of instances labeled as 'fresh' after resampled:", num_fresh_resampled)

```

Number of instances labeled as 'rotten' after resampled: 3812

Number of instances labeled as 'fresh' after resampled: 3812

## Logistic Regression Classifier

```

In [ ]: logisticModel = LogisticRegression()
logisticModel.fit(x_train, y_train)

```

```

Out[ ]: LogisticRegression
LogisticRegression()

```

```

In [ ]: logisticModelAccuracy = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))
logisticModelPrecision = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))
logisticModelSensitivity = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))
logisticModelF1Score = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))
logisticModelMCCScore = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))
logisticModelKappaCoeff = np.mean(cross_val_score(logisticModel, x_train, y_train, cv=5))

```

```

In [ ]: print("Train Accuracy: ", logisticModelAccuracy*100)
print("Train Precision: ", logisticModelPrecision*100)
print("Train Sensitivity: ", logisticModelSensitivity*100)
print("Train F1 Score: ", logisticModelF1Score*100)
print("Train MCC Score: ", logisticModelMCCScore*100)
print("Train Kappa Coefficient: ", logisticModelKappaCoeff*100)

```

Train Accuracy: 7371.494004425915

Train Precision: 7394.028016863263

Train Sensitivity: 7337.3318569553785

Train F1 Score: 7360.645739684533

Train MCC Score: 4749.160972824706

Train Kappa Coefficient: 4742.882443766669

```
In [ ]: y_pred_logistic = logisticModel.predict(x_test)
```

```
In [ ]: mse_logistic = mean_squared_error(y_test, y_pred_logistic)
rmse_logistic = np.sqrt(mse_logistic)
mae_logistic = mean_absolute_error(y_test, y_pred_logistic)

print(f"MSE: {mse_logistic}")
print(f"RMSE: {rmse_logistic}")
print(f"MAE: {mae_logistic}")
```

```
MSE: 0.2757679180887372
RMSE: 0.5251360948256529
MAE: 0.2757679180887372
```

```
In [ ]: logisticModelTestAccuracy = accuracy_score(y_test, y_pred_logistic)*100
logisticModelTestPrecision = precision_score(y_test, y_pred_logistic)*100
logisticModelTestSenstivity = recall_score(y_test, y_pred_logistic)*100
logisticModelTestF1Score = f1_score(y_test, y_pred_logistic)*100
logisticModelTestMCCScore = matthews_corrcoef(y_test, y_pred_logistic)*10
logisticModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_logistic)*

print(f"Test Accuracy: {logisticModelTestAccuracy*100}")
print(f"Test Precision: {logisticModelPrecision*100}")
print(f"Test Senstivity: {logisticModelSenstivity*100}")
print(f"Test F1 Score: {logisticModelF1Score*100}")
print(f"Test MCC Score: {logisticModelMCCScore*100}")
print(f"Test Kappa Coefficient: {logisticModelKappaCoeff*100}")
```


```
Test Accuracy: 7242.320819112628
Test Precision: 7394.028016863263
Test Senstivity: 7337.3318569553785
Test F1 Score: 7360.645739684533
Test MCC Score: 4749.160972824706
Test Kappa Coefficient: 4742.882443766669
```

```
In [ ]: report_logistic = classification_report(y_test, y_pred_logistic)
print(report_logistic)
```

	precision	recall	f1-score	support
0	0.76	0.74	0.75	1622
1	0.69	0.70	0.69	1308
accuracy			0.72	2930
macro avg	0.72	0.72	0.72	2930
weighted avg	0.72	0.72	0.72	2930

```
In [ ]: cm_logistic = confusion_matrix(y_test, y_pred_logistic)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_logistic)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



## Support Vector Machine Classifier (SVM)

```
In [ ]: svmModel = SVC()
        svmModel.fit(x_train, y_train)
```

```
Out [ ]: SVC
         SVC()
```

```
In [ ]: svmModelAccuracy = np.mean(cross_val_score(svmModel, x_train, y_train, sc
svmModelPrecision = np.mean(cross_val_score(svmModel, x_train, y_train, s
svmModelSensitivity = np.mean(cross_val_score(svmModel, x_train, y_train,
svmModelF1Score = np.mean(cross_val_score(svmModel, x_train, y_train, sco
svmModelMCCScore = np.mean(cross_val_score(svmModel, x_train, y_train, sc
svmModelKappaCoeff = np.mean(cross_val_score(svmModel, x_train, y_train,
```

```
In [ ]: print("Train Accuracy: ", svmModelAccuracy*100)
        print("Train Precision: ", svmModelPrecision*100)
        print("Train Sensitivity: ", svmModelSensitivity*100)
        print("Train F1 Score: ", svmModelF1Score*100)
        print("Train MCC Score: ", svmModelMCCScore*100)
        print("Train Kappa Coefficient: ", svmModelKappaCoeff*100)
```

Train Accuracy: 8549.297514281303  
 Train Precision: 8596.243453940791  
 Train Sensitivity: 8494.156003937009  
 Train F1 Score: 8539.495363566963  
 Train MCC Score: 7107.630400589809  
 Train Kappa Coefficient: 7098.558190785962

```
In [ ]: y_pred_svm = svmModel.predict(x_test)
```

```
In [ ]: mse_svm = mean_squared_error(y_test, y_pred_svm)
rmse_svm = np.sqrt(mse_svm)
mae_svm = mean_absolute_error(y_test, y_pred_svm)

print(f"MSE: {mse_svm}")
print(f"RMSE: {rmse_svm}")
print(f"MAE: {mae_svm}")
```

MSE: 0.15563139931740613  
 RMSE: 0.39450145667336406  
 MAE: 0.15563139931740613

```
In [ ]: svmModelTestAccuracy = accuracy_score(y_test, y_pred_svm)*100
svmModelTestPrecision = precision_score(y_test, y_pred_svm)*100
svmModelTestSensitivity = recall_score(y_test, y_pred_svm)*100
svmModelTestF1Score = f1_score(y_test, y_pred_svm)*100
svmModelTestMCCScore = matthews_corrcoef(y_test, y_pred_svm)*100
svmModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_svm)*100

print(f"Test Accuracy: {svmModelTestAccuracy*100}")
print(f"Test Precision: {svmModelTestPrecision*100}")
print(f"Test Sensitivity: {svmModelTestSensitivity*100}")
print(f"Test F1 Score: {svmModelTestF1Score*100}")
print(f"Test MCC Score: {svmModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {svmModelTestKappaCoeff*100}")
```


Test Accuracy: 8443.68600682594  
 Test Precision: 8297.213622291023  
 Test Sensitivity: 8195.718654434251  
 Test F1 Score: 8246.153846153848  
 Test MCC Score: 6847.895010536217  
 Test Kappa Coefficient: 6847.476055465998

```
In [ ]: report_svm = classification_report(y_test, y_pred_svm)
print(report_svm)
```

	precision	recall	f1-score	support
0	0.86	0.86	0.86	1622
1	0.83	0.82	0.82	1308
accuracy			0.84	2930
macro avg	0.84	0.84	0.84	2930
weighted avg	0.84	0.84	0.84	2930

```
In [ ]: cm_svm = confusion_matrix(y_test, y_pred_svm)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_svm)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



## Decision Tree Classifier

```
In [ ]: decisionTreeModel = DecisionTreeClassifier()
decisionTreeModel.fit(x_train, y_train)
```

```
Out [ ]: ▼ DecisionTreeClassifier ⓘ ⓘ
DecisionTreeClassifier()
```

```
In [ ]: decisionTreeModelAccuracy = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
decisionTreeModelPrecision = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
decisionTreeModelSensitivity = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
decisionTreeModelF1Score = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
decisionTreeModelMCCScore = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
decisionTreeModelKappaCoeff = np.mean(cross_val_score(decisionTreeModel, x_train, y_train, cv=5))
```

```
In [ ]: print("Train Accuracy: ", decisionTreeModelAccuracy*100)
print("Train Precision: ", decisionTreeModelPrecision*100)
```



```
print("Train Sensitivity: ", decisionTreeModelSensitivity*100)
print("Train F1 Score: ", decisionTreeModelF1Score*100)
print("Train MCC Score: ", decisionTreeModelMCCScore*100)
print("Train Kappa Coefficient: ", decisionTreeModelKappaCoeff*100)
```

```
Train Accuracy: 8641.145592095105
Train Precision: 8672.524078859797
Train Sensitivity: 8580.790682414698
Train F1 Score: 8635.653483094515
Train MCC Score: 7281.0656401081405
Train Kappa Coefficient: 7316.385459108208
```

```
In [ ]: y_pred_decision_tree = decisionTreeModel.predict(x_test)
```

```
In [ ]: mse_decision_tree = mean_squared_error(y_test, y_pred_decision_tree)
rmse_decision_tree = np.sqrt(mse_decision_tree)
mae_decision_tree = mean_absolute_error(y_test, y_pred_decision_tree)

print(f"MSE: {mse_decision_tree}")
print(f"RMSE: {rmse_decision_tree}")
print(f"MAE: {mae_decision_tree}")
```

```
MSE: 0.14948805460750852
RMSE: 0.38663685107282325
MAE: 0.14948805460750852
```

```
In [ ]: decisionTreeModelTestAccuracy = accuracy_score(y_test, y_pred_decision_tr
decisionTreeModelTestPrecision = precision_score(y_test, y_pred_decision_
decisionTreeModelTestSensitivity = recall_score(y_test, y_pred_decision_tr
decisionTreeModelTestF1Score = f1_score(y_test, y_pred_decision_tree)*100
decisionTreeModelTestMCCScore = matthews_corrcoef(y_test, y_pred_decision
decisionTreeModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_decisi

print(f"Test Accuracy: {decisionTreeModelTestAccuracy*100}")
print(f"Test Precision: {decisionTreeModelTestPrecision*100}")
print(f"Test Sensitivity: {decisionTreeModelTestSensitivity*100}")
print(f"Test F1 Score: {decisionTreeModelTestF1Score*100}")
print(f"Test MCC Score: {decisionTreeModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {decisionTreeModelTestKappaCoeff*100}")
```


```
Test Accuracy: 8505.119453924914
Test Precision: 8377.329192546584
Test Sensitivity: 8249.235474006116
Test F1 Score: 8312.788906009246
Test MCC Score: 6971.686987690813
Test Kappa Coefficient: 6971.020129944921
```

```
In [ ]: report_decision_tree = classification_report(y_test, y_pred_decision_tree)
print(report_decision_tree)
```

	precision	recall	f1-score	support
0	0.86	0.87	0.87	1622
1	0.84	0.82	0.83	1308
accuracy			0.85	2930
macro avg	0.85	0.85	0.85	2930
weighted avg	0.85	0.85	0.85	2930

```
In [ ]: cm_decision_tree = confusion_matrix(y_test, y_pred_decision_tree)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_decision_tree)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



## Random Forest Classifier

```
In [ ]: randomForestModel = RandomForestClassifier()
randomForestModel.fit(x_train, y_train)
```

```
Out [ ]: ▼ RandomForestClassifier ⓘ ⓘ
RandomForestClassifier()
```

```
In [ ]: randomForestModelAccuracy = np.mean(cross_val_score(randomForestModel, x_
randomForestModelPrecision = np.mean(cross_val_score(randomForestModel, x_
randomForestModelSensitivity = np.mean(cross_val_score(randomForestModel,
randomForestModelF1Score = np.mean(cross_val_score(randomForestModel, x_t
randomForestModelMCCScore = np.mean(cross_val_score(randomForestModel, x_
randomForestModelKappaCoeff = np.mean(cross_val_score(randomForestModel,
```

```
In [ ]: print("Accuracy: ", randomForestModelAccuracy*100)
print("Precision: ", randomForestModelPrecision*100)
print("Sensitivity: ", randomForestModelSensitivity*100)
print("F1 Score: ", randomForestModelF1Score*100)
print("MCC Score: ", randomForestModelMCCScore*100)
print("Kappa Coefficient: ", randomForestModelKappaCoeff*100)
```

```
Accuracy: 9154.037362976685
Precision: 9174.92087057908
Sensitivity: 9176.201607611547
F1 Score: 9152.793694569416
MCC Score: 8326.089987101403
Kappa Coefficient: 8318.549968458121
```

```
In [ ]: y_pred_random_forest = randomForestModel.predict(x_test)
```

```
In [ ]: mse_random_forest = mean_squared_error(y_test, y_pred_random_forest)
rmse_random_forest = np.sqrt(mse_random_forest)
mae_random_forest = mean_absolute_error(y_test, y_pred_random_forest)

print(f"MSE: {mse_random_forest}")
print(f"RMSE: {rmse_random_forest}")
print(f"MAE: {mae_random_forest}")
```

```
MSE: 0.1
RMSE: 0.31622776601683794
MAE: 0.1
```

```
In [ ]: randomForestModelTestAccuracy = accuracy_score(y_test, y_pred_random_fore
randomForestModelTestPrecision = precision_score(y_test, y_pred_random_fore
randomForestModelTestSensitivity = recall_score(y_test, y_pred_random_fore
randomForestTestModelF1Score = f1_score(y_test, y_pred_random_forest)*100
randomForestModelTestMCCScore = matthews_corrcoef(y_test, y_pred_random_f
randomForestModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_random

print(f"Test Accuracy: {randomForestModelTestAccuracy*100}")
print(f"Test Precision: {randomForestModelTestPrecision*100}")
print(f"Test Sensitivity: {randomForestModelTestSensitivity*100}")
print(f"Test F1 Score: {randomForestTestModelF1Score*100}")
print(f"Test MCC Score: {randomForestModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {randomForestModelTestKappaCoeff*100}")
```


Test Accuracy: 9000.0  
Test Precision: 8980.392156862747  
Test Sensitivity: 8753.822629969418  
Test F1 Score: 8865.66008517228  
Test MCC Score: 7973.891241731465  
Test Kappa Coefficient: 7971.810555212993

```
In [ ]: report_random_forest = classification_report(y_test, y_pred_random_forest)
print(report_random_forest)
```

	precision	recall	f1-score	support
0	0.90	0.92	0.91	1622
1	0.90	0.88	0.89	1308
accuracy			0.90	2930
macro avg	0.90	0.90	0.90	2930
weighted avg	0.90	0.90	0.90	2930

```
In [ ]: cm_random_forest = confusion_matrix(y_test, y_pred_random_forest)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_random_forest)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



# K Neighbors Classifier (KNN)

```
In [ ]: KNNModel = KNeighborsClassifier()
        KNNModel.fit(x_train, y_train)
```

```
Out [ ]: ▼ KNeighborsClassifier ⓘ ?
        KNeighborsClassifier()
```

```
In [ ]: KNNModelAccuracy = np.mean(cross_val_score(KNNModel, x_train, y_train, sc
KNNModelPrecision = np.mean(cross_val_score(KNNModel, x_train, y_train, s
KNNModelSensitivity = np.mean(cross_val_score(KNNModel, x_train, y_train, sc
KNNModelF1Score = np.mean(cross_val_score(KNNModel, x_train, y_train, sco
KNNModelMCCScore = np.mean(cross_val_score(KNNModel, x_train, y_train, sc
KNNModelKappaCoeff = np.mean(cross_val_score(KNNModel, x_train, y_train,
```

```
In [ ]: print("Train Accuracy: ", KNNModelAccuracy*100)
        print("Train Precision: ", KNNModelPrecision*100)
        print("Train Sensitivity: ", KNNModelSensitivity*100)
        print("Train F1 Score: ", KNNModelF1Score*100)
        print("Train MCC Score: ", KNNModelMCCScore*100)
        print("Train Kappa Coefficient: ", KNNModelKappaCoeff*100)
```

```
Train Accuracy: 8986.166435077967
Train Precision: 8881.664829027555
Train Sensitivity: 9131.6437007874
Train F1 Score: 8998.761372177041
Train MCC Score: 7987.5253561829495
Train Kappa Coefficient: 7972.270489286548
```

```
In [ ]: y_pred_knn = KNNModel.predict(x_test)
```

```
In [ ]: mse_knn = mean_squared_error(y_test, y_pred_knn)
        rmse_knn = np.sqrt(mse_knn)
        mae_knn = mean_absolute_error(y_test, y_pred_knn)

        print(f"MSE: {mse_knn}")
        print(f"RMSE: {rmse_knn}")
        print(f"MAE: {mae_knn}")
```

```
MSE: 0.12627986348122866
RMSE: 0.3553587813481308
MAE: 0.12627986348122866
```

```
In [ ]: KNNModelTestAccuracy = accuracy_score(y_test, y_pred_knn)*100
        KNNModelTestPrecision = precision_score(y_test, y_pred_knn)*100
        KNNModelTestSensitivity = recall_score(y_test, y_pred_knn)*100
        KNNModelTestF1Score = f1_score(y_test, y_pred_knn)*100
        KNNModelTestMCCScore = matthews_corrcoef(y_test, y_pred_knn)*100
        KNNModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_knn)*100
```

```

print(f"Test Accuracy: {KNNModelTestAccuracy*100}")
print(f"Test Precision: {KNNModelTestPrecision*100}")
print(f"Test Sensitivity: {KNNModelTestSensitivity*100}")
print(f"Test F1 Score: {KNNModelTestF1Score*100}")
print(f"Test MCC Score: {KNNModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {KNNModelTestKappaCoeff*100}")

```

```

Test Accuracy: 8737.201365187713
Test Precision: 8510.479041916167
Test Sensitivity: 8692.660550458715
Test F1 Score: 8600.605143721634
Test MCC Score: 7451.729928801278
Test Kappa Coefficient: 7450.342713826899

```

```

In [ ]: report_knn = classification_report(y_test, y_pred_knn)
        print(report_knn)

```


	precision	recall	f1-score	support
0	0.89	0.88	0.88	1622
1	0.85	0.87	0.86	1308
accuracy			0.87	2930
macro avg	0.87	0.87	0.87	2930
weighted avg	0.87	0.87	0.87	2930

```

In [ ]: cm_knn = confusion_matrix(y_test, y_pred_knn)

        disp = ConfusionMatrixDisplay(confusion_matrix=cm_knn)
        disp.plot(cmap=plt.cm.Blues)
        plt.show()

```



## Naive Bayes Classifier

```
In [ ]: naiveBayesModel = GaussianNB()
naiveBayesModel.fit(x_train, y_train)
```

```
Out [ ]: GaussianNB
GaussianNB()
```

```
In [ ]: naiveBayesModelAccuracy = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
naiveBayesModelPrecision = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
naiveBayesModelSensitivity = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
naiveBayesModelF1Score = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
naiveBayesModelMCCScore = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
naiveBayesModelKappaCoeff = np.mean(cross_val_score(naiveBayesModel, x_train, y_train, cv=5))
```

```
In [ ]: print("Train Accuracy: ", naiveBayesModelAccuracy*100)
print("Train Precision: ", naiveBayesModelPrecision*100)
print("Train Sensitivity: ", naiveBayesModelSensitivity*100)
print("Train F1 Score: ", naiveBayesModelF1Score*100)
print("Train MCC Score: ", naiveBayesModelMCCScore*100)
print("Train Kappa Coefficient: ", naiveBayesModelKappaCoeff*100)
```

```

Train Accuracy: 6537.296073284958
Train Precision: 6479.190442912921
Train Sensitivity: 6739.173228346456
Train F1 Score: 6601.154551828718
Train MCC Score: 3082.412406496794
Train Kappa Coefficient: 3074.4519545851017

```

```
In [ ]: y_pred_naive_bayes = naiveBayesModel.predict(x_test)
```

```
In [ ]: mse_naive_bayes = mean_squared_error(y_test, y_pred_naive_bayes)
rmse_naive_bayes = np.sqrt(mse_naive_bayes)
mae_naive_bayes = mean_absolute_error(y_test, y_pred_naive_bayes)

print(f"MSE: {mse_naive_bayes}")
print(f"RMSE: {rmse_naive_bayes}")
print(f"MAE: {mae_naive_bayes}")
```

```

MSE: 0.37679180887372016
RMSE: 0.6138336980597596
MAE: 0.37679180887372016

```

```
In [ ]: naiveBayesModelTestAccuracy = accuracy_score(y_test, y_pred_naive_bayes)*
naiveBayesModelTestPrecision = precision_score(y_test, y_pred_naive_bayes)
naiveBayesModelTestSensitivity = recall_score(y_test, y_pred_naive_bayes)*
naiveBayesModelTestF1Score = f1_score(y_test, y_pred_naive_bayes)*100
naiveBayesModelTestMCCScore = matthews_corrcoef(y_test, y_pred_naive_bayes)
naiveBayesModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_naive_bayes)

print(f"Test Accuracy: {naiveBayesModelTestAccuracy*100}")
print(f"Test Precision: {naiveBayesModelTestPrecision*100}")
print(f"Test Sensitivity: {naiveBayesModelTestSensitivity*100}")
print(f"Test F1 Score: {naiveBayesModelTestF1Score*100}")
print(f"Test MCC Score: {naiveBayesModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {naiveBayesModelTestKappaCoeff*100}")
```

```

Test Accuracy: 6232.081911262799
Test Precision: 5698.630136986301
Test Sensitivity: 6360.856269113149
Test F1 Score: 6011.560693641619
Test MCC Score: 2474.772762744032
Test Kappa Coefficient: 2461.4065180102907

```


```
In [ ]: report_naive_bayes = classification_report(y_test, y_pred_naive_bayes)
print(report_naive_bayes)
```

	precision	recall	f1-score	support
0	0.68	0.61	0.64	1622
1	0.57	0.64	0.60	1308
accuracy			0.62	2930
macro avg	0.62	0.62	0.62	2930
weighted avg	0.63	0.62	0.62	2930



```
In [ ]: cm_naive_bayes = confusion_matrix(y_test, y_pred_naive_bayes)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_naive_bayes)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



## XGBoost

```
In [ ]: xgboostModel = XGBClassifier()
xgboostModel.fit(x_train, y_train)
```

```
Out [ ]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_
rounds=None,
              enable_categorical=False, eval_metric=None, feature
_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, m
ax_bin=None,
```

```
In [ ]: xgboostModelAccuracy = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
xgboostModelPrecision = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
xgboostModelSensitivity = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
xgboostModelF1Score = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
xgboostModelMCCScore = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
xgboostModelKappaCoeff = np.mean(cross_val_score(xgboostModel, x_train, y_train, cv=5))
```

```
In [ ]: print("Train Accuracy: ", xgboostModelAccuracy*100)
print("Train Precision: ", xgboostModelPrecision*100)
print("Train Sensitivity: ", xgboostModelSensitivity*100)
print("Train F1 Score: ", xgboostModelF1Score*100)
print("Train MCC Score: ", xgboostModelMCCScore*100)
print("Train Kappa Coefficient: ", xgboostModelKappaCoeff*100)
```

```
Train Accuracy: 90.16314137203437
Train Precision: 89.71703535512197
Train Sensitivity: 90.79191272965876
Train F1 Score: 90.21351739753547
Train MCC Score: 80.40349648550983
Train Kappa Coefficient: 80.32590553723433
```

```
In [ ]: y_pred_xgboost = xgboostModel.predict(x_test)
```

```
In [ ]: mse_xgboost = mean_squared_error(y_test, y_pred_xgboost)
rmse_xgboost = np.sqrt(mse_xgboost)
mae_xgboost = mean_absolute_error(y_test, y_pred_xgboost)

print(f"MSE: {mse_xgboost}")
print(f"RMSE: {rmse_xgboost}")
print(f"MAE: {mae_xgboost}")
```

```
MSE: 0.10273037542662115
RMSE: 0.3205157959081286
MAE: 0.10273037542662115
```

```
In [ ]: xgboostModelTestAccuracy = accuracy_score(y_test, y_pred_xgboost)*100
xgboostModelTestPrecision = precision_score(y_test, y_pred_xgboost)*100
xgboostModelTestSensitivity = recall_score(y_test, y_pred_xgboost)*100
xgboostModelTestF1Score = f1_score(y_test, y_pred_xgboost)*100
xgboostModelTestMCCScore = matthews_corrcoef(y_test, y_pred_xgboost)*100
xgboostModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_xgboost)*100

print(f"Test Accuracy: {xgboostModelTestAccuracy*100}")
print(f"Test Precision: {xgboostModelTestPrecision*100}")
print(f"Test Sensitivity: {xgboostModelTestSensitivity*100}")
print(f"Test F1 Score: {xgboostModelTestF1Score*100}")
print(f"Test MCC Score: {xgboostModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {xgboostModelTestKappaCoeff*100}")
```


Test Accuracy: 8972.696245733789  
Test Precision: 8906.128782001551  
Test Sensitivity: 8776.758409785933  
Test F1 Score: 8840.970350404312  
Test MCC Score: 7919.278566381003  
Test Kappa Coefficient: 7918.5950298146

```
In [ ]: report_xgboost = classification_report(y_test, y_pred_xgboost)
print(report_xgboost)
```

	precision	recall	f1-score	support
0	0.90	0.91	0.91	1622
1	0.89	0.88	0.88	1308
accuracy			0.90	2930
macro avg	0.90	0.90	0.90	2930
weighted avg	0.90	0.90	0.90	2930

```
In [ ]: cm_xgboost = confusion_matrix(y_test, y_pred_xgboost)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_xgboost)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



# AdaBoost

```
In [ ]: adaboostModel = AdaBoostClassifier()
        adaboostModel.fit(x_train, y_train)
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

```
Out[ ]: ▼ AdaBoostClassifier ⓘ ?
        AdaBoostClassifier()
```

```
In [ ]: adaboostModelAccuracy = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
        adaboostModelPrecision = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
        adaboostModelSensitivity = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
        adaboostModelF1Score = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
        adaboostModelLMCCScore = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
        adaboostModelKappaCoeff = np.mean(cross_val_score(adaboostModel, x_train, y_train, cv=5))
```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
warnings.warn(

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(

```

```
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME al
```



```
algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
        warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
            warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
                warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
                    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
```

[illegible]



```
earn/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
```

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(

```

```
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME al
```

```
algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
```

[illegible]



```

earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
  warnings.warn(

```

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(

```

```
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME al
```



```
algorithm to circumvent this warning.  
    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
        warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
            warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
                warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.  
                    warnings.warn(  
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
```

[illegible]

```
earn/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/skl
earn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorit
m (the default) is deprecated and will be removed in 1.6. Use the SAMME al
gorithm to circumvent this warning.
```

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
  warnings.warn(

```

```
warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
```

```
In [ ]: print("Train Accuracy: ", adaboostModelAccuracy*100)
print("Train Precision: ", adaboostModelPrecision*100)
print("Train Sensitivity: ", adaboostModelSensitivity*100)
print("Train F1 Score: ", adaboostModelF1Score*100)
print("Train MCC Score: ", adaboostModelMCCScore*100)
print("Train Kappa Coefficient: ", adaboostModelKappaCoeff*100)
```

```
Train Accuracy: 7548.546137615149
Train Precision: 7512.617461709035
Train Sensitivity: 7636.380413385825
Train F1 Score: 7568.58310714984
Train MCC Score: 5105.205235782601
Train Kappa Coefficient: 5096.984060993945
```

```
In [ ]: y_pred_adaboost = adaboostModel.predict(x_test)
```

```
In [ ]: mse_adaboost = mean_squared_error(y_test, y_pred_adaboost)
rmse_adaboost = np.sqrt(mse_adaboost)
mae_adaboost = mean_absolute_error(y_test, y_pred_adaboost)

print(f"MSE: {mse_adaboost}")
print(f"RMSE: {rmse_adaboost}")
print(f"MAE: {mae_adaboost}")
```

```
MSE: 0.24914675767918087
RMSE: 0.49914602841170724
MAE: 0.24914675767918087
```

```
In [ ]: adaboostModelTestAccuracy = accuracy_score(y_test, y_pred_adaboost)*100
adaboostModelTestPrecision = precision_score(y_test, y_pred_adaboost)*100
adaboostModelTestSensitivity = recall_score(y_test, y_pred_adaboost)*100
adaboostModelTestF1Score = f1_score(y_test, y_pred_adaboost)*100
adaboostModelTestMCCScore = matthews_corrcoef(y_test, y_pred_adaboost)*100
adaboostModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_adaboost)*100

print(f"Test Accuracy: {adaboostModelTestAccuracy*100}")
print(f"Test Precision: {adaboostModelTestPrecision*100}")
print(f"Test Sensitivity: {adaboostModelTestSensitivity*100}")
print(f"Test F1 Score: {adaboostModelTestF1Score*100}")
print(f"Test MCC Score: {adaboostModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {adaboostModelTestKappaCoeff*100}")
```


Test Accuracy: 7508.53242320819  
Test Precision: 7118.768328445748  
Test Sensitivity: 7423.547400611621  
Test F1 Score: 7267.964071856288  
Test MCC Score: 4983.672062170667  
Test Kappa Coefficient: 4979.975290654009

```
In [ ]: report_adaboost = classification_report(y_test, y_pred_adaboost)
print(report_adaboost)
```

	precision	recall	f1-score	support
0	0.78	0.76	0.77	1622
1	0.71	0.74	0.73	1308
accuracy			0.75	2930
macro avg	0.75	0.75	0.75	2930
weighted avg	0.75	0.75	0.75	2930

```
In [ ]: cm_adaboost = confusion_matrix(y_test, y_pred_adaboost)

disp = ConfusionMatrixDisplay(confusion_matrix=cm_adaboost)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```





# Model Comparison

## Metrics Comparison on Train Data

### Classification Performance Metrics

```
In [ ]: arr_train = np.array(["Logistic Regression", logisticModelAccuracy, logi
arr_train = np.insert(arr_train, 0, ["SVM", svmModelAccuracy, svmModelPre
arr_train = np.insert(arr_train, 0, ["Decision Tree", decisionTreeModelAc
arr_train = np.insert(arr_train, 0, ["Random Forest", randomForestModelAc
arr_train = np.insert(arr_train, 0, ["KNN", KNNModelAccuracy, KNNModelPre
arr_train = np.insert(arr_train, 0, ["Naive Bayes", naiveBayesModelAccura
arr_train = np.insert(arr_train, 0, ["XGBoost", xgboostModelAccuracy, xgb
arr_train = np.insert(arr_train, 0, ["AdaBoost", adaboostModelAccuracy, a
```

```
In [ ]: arr_train = pd.DataFrame(arr_train, columns=['Model', 'Accuracy', 'Precis
```

```
In [ ]: arr_train
```

Out[ ]:

	Model	Accuracy	Precision	Sensitivity	
0	AdaBoost	75.48546137615149	75.12617461709034	76.36380413385825	75.6
1	XGBoost	90.16314137203437	89.71703535512196	90.79191272965876	90.21
2	Naive Bayes	65.37296073284958	64.7919044291292	67.39173228346456	66.0
3	KNN	89.86166435077966	88.81664829027555	91.316437007874	89.98
4	Random Forest	91.54037362976685	91.7492087057908	91.76201607611547	91.52
5	Decision Tree	86.41145592095104	86.72524078859796	85.80790682414698	86.35
6	SVM	85.49297514281302	85.96243453940791	84.94156003937009	85.39
7	Logistic Regression	73.71494004425915	73.94028016863263	73.37331856955379	73.60

\*\* Random Forest and XG Boost classifier are very close to each other. Hence I am going to use both of them with different hyperameters to see which one performs better.

## Metrics Comparison on Test Data

### Classification Performance Metrics

```
In [ ]: arr_test = np.array(["Logistic Regression", logisticModelTestAccuracy, 1
arr_test = np.insert(arr_test, 0, ["SVM", svmModelTestAccuracy, svmModelT
arr_test = np.insert(arr_test, 0, ["Decision Tree", decisionTreeModelTest
arr_test = np.insert(arr_test, 0, ["Random Forest", randomForestModelTest
arr_test = np.insert(arr_test, 0, ["KNN", KNNModelTestAccuracy, KNNModelT
arr_test = np.insert(arr_test, 0, ["Naive Bayes", naiveBayesModelTestAccu
arr_test = np.insert(arr_test, 0, ["XGBoost", xgboostModelTestAccuracy, x
arr_test = np.insert(arr_test, 0, ["AdaBoost", adaboostModelTestAccuracy,
```

```
In [ ]: arr_test = pd.DataFrame(arr_test, columns=['Model', 'Accuracy', 'Precisio
```

```
In [ ]: arr_test
```

```
Out[ ]:
```

	Model	Accuracy	Precision	Sensitivity	
0	AdaBoost	75.0853242320819	71.18768328445748	74.2354740061162	72.
1	XGBoost	89.72696245733789	89.06128782001551	87.76758409785933	88.
2	Naive Bayes	62.32081911262799	56.986301369863014	63.608562691131496	60.1
3	KNN	87.37201365187714	85.10479041916167	86.92660550458714	86.
4	Random Forest	90.0	89.80392156862746	87.53822629969419	88.
5	Decision Tree	85.05119453924914	83.77329192546584	82.49235474006116	83.
6	SVM	84.43686006825939	82.97213622291022	81.9571865443425	82.
7	Logistic Regression	72.42320819112628	68.65671641791045	70.33639143730886	69.4

## Regression Performance Metrics

```
In [ ]: arr_regression = np.array(["Logistic Regression", mse_logistic, rmse_log
arr_regression = np.insert(arr_regression, 0, ["SVM", mse_svm, rmse_svm,
arr_regression = np.insert(arr_regression, 0, ["Decision Tree", mse_decis
arr_regression = np.insert(arr_regression, 0, ["Random Forest", mse_rando
arr_regression = np.insert(arr_regression, 0, ["KNN", mse_knn, rmse_knn,
arr_regression = np.insert(arr_regression, 0, ["Naive Bayes", mse_naive_b
arr_regression = np.insert(arr_regression, 0, ["XGBoost", mse_xgboost, rm
arr_regression = np.insert(arr_regression, 0, ["AdaBoost", mse_adaboost,
```

```
In [ ]: arr_regression = pd.DataFrame(arr_regression, columns=['Model', 'MSE', 'R
```

```
In [ ]: arr_regression
```



Out[ ]:

	Model	MSE		RMSE	MAE
0	AdaBoost	0.24914675767918087	0.49914602841170724	0.24914675767918087	
1	XGBoost	0.10273037542662115	0.3205157959081286	0.10273037542662115	
2	Naive Bayes	0.37679180887372016	0.6138336980597596	0.37679180887372016	
3	KNN	0.12627986348122866	0.3553587813481308	0.12627986348122866	
4	Random Forest	0.1	0.31622776601683794		0.1
5	Decision Tree	0.14948805460750852	0.38663685107282325	0.14948805460750852	
6	SVM	0.15563139931740613	0.39450145667336406	0.15563139931740613	
7	Logistic Regression	0.2757679180887372	0.5251360948256529	0.2757679180887372	

\*\* The lower the value, the better the performance of the model

## Final Comparison: Random Forest vs XGBoost using GridSearch

```
In [ ]: # Random Forest Hyperparameter Grid
random_forest_param_grid = {
    'n_estimators': [100, 200, 300, 400, 500],           # Number of trees
    'max_depth': [None, 10, 20, 30, 40],                 # Maximum depth
    'min_samples_split': [2, 5, 10, 15, 20],             # Minimum number of samples to split
    'min_samples_leaf': [1, 2, 4, 6, 8],                # Minimum number of samples in leaf
    'max_features': ['auto', 'sqrt', 'log2', 0.2, 0.4]   # Number of features to use
}

# Define Scoring Metrics
scorers = {
    'accuracy': make_scorer(accuracy_score),
    'precision': make_scorer(precision_score),
    'sensitivity': make_scorer(recall_score),
    'f1': make_scorer(f1_score),
    'mcc': make_scorer(matthews_corrcoef),
    'kappa': make_scorer(cohen_kappa_score)
}

# Random Forest Grid Search
random_forest_model = RandomForestClassifier(random_state=42)
grid_search_rf = GridSearchCV(
    estimator=random_forest_model,
    param_grid=random_forest_param_grid,
    scoring=scorers,
```

```

    refit='f1',
    cv=5,
    return_train_score=True
)
grid_search_rf.fit(x_train, y_train)

```

/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model\_selection/\_validation.py:547: FitFailedWarning:  
3125 fits failed out of a total of 15625.  
The score on these train-test partitions for these parameters will be set to nan.  
If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

```

-----
3125 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_validation.py", line 895, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/base.py", line 1467, in wrapper
    estimator._validate_params()
  File "/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/base.py", line 666, in _validate_params
    validate_parameter_constraints(
  File "/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/utils/_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test scores are non-finite: [          nan          nan          nan ... 0.87618304 0.8743468 0.87605232]
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the train scores are non-finite: [          nan          nan          nan ... 0.93510634 0.93441769 0.9349424 ]
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test scores are non-finite: [          nan          nan          nan ... 0.87336709 0.87019617 0.87268716]
    warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sk

```

```

earn/model_selection/_search.py:1051: UserWarning: One or more of the train
scores are non-finite: [          nan          nan          nan ... 0.93382883
0.93317052 0.93369425]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test
scores are non-finite: [          nan          nan          nan ... 0.88012164 0.
88012198 0.88064623]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the train
scores are non-finite: [          nan          nan          nan ... 0.93658208
0.93586064 0.93638531]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test
scores are non-finite: [          nan          nan          nan ... 0.8764682 0.
87486064 0.87638758]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the train
scores are non-finite: [          nan          nan          nan ... 0.93520001
0.93450964 0.9350333 ]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test
scores are non-finite: [          nan          nan          nan ... 0.75288649 0.
74926091 0.75263509]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the train
scores are non-finite: [          nan          nan          nan ... 0.87022292
0.86884652 0.86989694]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the test
scores are non-finite: [          nan          nan          nan ... 0.75236645 0.
74869418 0.75210511]
  warnings.warn(
/Users/brahderlau/anaconda3/envs/capstone/lib/python3.10/site-packages/sklearn/model_selection/_search.py:1051: UserWarning: One or more of the train
scores are non-finite: [          nan          nan          nan ... 0.87021269
0.86883538 0.86988479]
  warnings.warn(

```

```

Out[ ]:  ► GridSearchCV ⓘ ⓘ
          ► estimator: RandomForestClassifier
            ► RandomForestClassifier ⓘ

```

```

In [ ]: # Evaluate Random Forest Best Model

```

```

rf_best = grid_search_rf.best_estimator_

# Feature importance
importance_values_rf = rf_best.feature_importances_

sorted_indices_rf = importance_values_rf.argsort()[::-1]
sorted_importances_rf = importance_values_rf[sorted_indices_rf]
sorted_feature_names_rf = [feature_names[i] for i in sorted_indices_rf]

plt.figure(figsize=(10, 6))
plt.barh(range(len(sorted_importances_rf)), sorted_importances_rf, tick_l
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Random Forest Feature Importances')
plt.gca().invert_yaxis()
plt.show()

rf_saved_name = "best_rf_model_" + str(grid_search_rf.best_params_) + ".p
with open(rf_saved_name, 'wb') as model_file:
    pickle.dump(grid_search_rf, model_file)

print("Best Random Forest Model saved successfully.")

y_pred_rf = rf_best.predict(x_test)
print("\nRandom Forest Best Parameters:", grid_search_rf.best_params_)

finalRfModelTestAccuracy = accuracy_score(y_test, y_pred_rf)*100
finalRfModelTestPrecision = precision_score(y_test, y_pred_rf)*100
finalRfModelTestSenstivity = recall_score(y_test, y_pred_rf)*100
finalRfModelTestF1Score = f1_score(y_test, y_pred_rf)*100
finalRfModelTestMCCScore = matthews_corrcoef(y_test, y_pred_rf)*100
finalRfModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_rf)*100

print(f"\nTest Accuracy: {finalRfModelTestAccuracy*100}")
print(f"Test Precision: {finalRfModelTestPrecision*100}")
print(f"Test Senstivity: {finalRfModelTestSenstivity*100}")
print(f"Test F1 Score: {finalRfModelTestF1Score*100}")
print(f"Test MCC Score: {finalRfModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {finalRfModelTestKappaCoeff*100}")


print("\nRandom Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))
cm_rf = confusion_matrix(y_test, y_pred_rf)
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf)
disp_rf.plot(cmap=plt.cm.Blues)
plt.show()

final_mse_rf = mean_squared_error(y_test, y_pred_rf)
final_rmse_rf = np.sqrt(final_mse_rf)
final_mae_rf = mean_absolute_error(y_test, y_pred_rf)

print(f"Final MSE: {final_mse_rf}")
print(f"Final RMSE: {final_rmse_rf}")

```

```
print(f"Final MAE: {final_mae_rf}")
```




Best Random Forest Model saved successfully.

Random Forest Best Parameters: {'max\_depth': None, 'max\_features': 'sqrt', 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}

Test Accuracy: 8948.805460750853  
Test Precision: 8937.007874015748  
Test Sensitivity: 8677.37003058104  
Test F1 Score: 8805.275407292474  
Test MCC Score: 7869.912365823544  
Test Kappa Coefficient: 7867.187241504815

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.90	0.92	0.91	1622
1	0.89	0.87	0.88	1308
accuracy			0.89	2930
macro avg	0.89	0.89	0.89	2930
weighted avg	0.89	0.89	0.89	2930



Final MSE: 0.10511945392491467


Final RMSE: 0.3242213039343878

Final MAE: 0.10511945392491467

```
In [ ]: # XGBoost Hyperparameter Grid
xgboost_param_grid = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [3, 5, 7, 10, 15],
    'learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3],
    'subsample': [0.7, 0.8, 0.85, 0.9, 1.0],
    'gamma': [0, 0.1, 0.2, 0.3, 0.5]
}

# XGBoost Grid Search
xgboost_model = XGBClassifier(random_state=42)
grid_search_xgb = GridSearchCV(
    estimator=xgboost_model,
    param_grid=xgboost_param_grid,
    scoring=scorers,
    refit='f1',
    cv=5,
    return_train_score=True
)
grid_search_xgb.fit(x_train, y_train)
```

Out [ ]:



In [ ]:

```

# Evaluate XGBoost Best Model
xgb_best = grid_search_xgb.best_estimator_

importance_values_xgb = xgb_best.feature_importances_

sorted_indices_xgb = importance_values_xgb.argsort()[::-1]
sorted_importances_xgb = importance_values_xgb[sorted_indices_xgb]
sorted_feature_names_xgb = [feature_names[i] for i in sorted_indices_xgb]

plt.figure(figsize=(10, 6))
plt.barh(range(len(sorted_importances_xgb)), sorted_importances_xgb, tick
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('XGB Feature Importances')
plt.gca().invert_yaxis()
plt.show()

xgb_saved_name = "best_xgb_model_" + str(grid_search_xgb.best_params_) +
with open(xgb_saved_name, 'wb') as model_file:
    pickle.dump(xgb_best, model_file)

print("Best XG Boost Model saved successfully.")

y_pred_xgb = xgb_best.predict(x_test)
print("XGBoost Best Parameters:", grid_search_xgb.best_params_)

finalXgbModelTestAccuracy = accuracy_score(y_test, y_pred_xgb)*100
finalXgbModelTestPrecision = precision_score(y_test, y_pred_xgb)*100
finalXgbModelTestSensitivity = recall_score(y_test, y_pred_xgb)*100
finalXgbModelTestF1Score = f1_score(y_test, y_pred_xgb)*100
finalXgbModelTestMCCScore = matthews_corrcoef(y_test, y_pred_xgb)*100
finalXgbModelTestKappaCoeff = cohen_kappa_score(y_test, y_pred_xgb)*100

print(f"\nTest Accuracy: {finalXgbModelTestAccuracy*100}")
print(f"Test Precision: {finalXgbModelTestPrecision*100}")
print(f"Test Sensitivity: {finalXgbModelTestSensitivity*100}")
print(f"Test F1 Score: {finalXgbModelTestF1Score*100}")
print(f"Test MCC Score: {finalXgbModelTestMCCScore*100}")
print(f"Test Kappa Coefficient: {finalXgbModelTestKappaCoeff*100}")


print("\nXGBoost Classification Report:")
print(classification_report(y_test, y_pred_xgb))
cm_xgb = confusion_matrix(y_test, y_pred_xgb)
disp_xgb = ConfusionMatrixDisplay(confusion_matrix=cm_xgb)
disp_xgb.plot(cmap=plt.cm.Blues)

```

```
plt.show()

final_mse_xgb = mean_squared_error(y_test, y_pred_xgb)
final_rmse_xgb= np.sqrt(final_mse_xgb)
final_mae_xgb= mean_absolute_error(y_test, y_pred_xgb)

print(f"Final MSE: {final_mse_xgb}")
print(f"Final RMSE: {final_rmse_xgb}")
print(f"Final MAE: {final_mae_xgb}")
```



Best XG Boost Model saved successfully.

XGBoost Best Parameters: {'gamma': 0.1, 'learning\_rate': 0.1, 'max\_depth': 15, 'n\_estimators': 300, 'subsample': 0.8}

Test Accuracy: 9058.0204778157

Test Precision: 8975.346687211095

Test Sensitivity: 8906.727828746178

Test F1 Score: 8940.90560245587


Test MCC Score: 8092.934480397588

Test Kappa Coefficient: 8092.741246958735

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.91	0.92	0.92	1622
1	0.90	0.89	0.89	1308
accuracy			0.91	2930
macro avg	0.90	0.90	0.90	2930
weighted avg	0.91	0.91	0.91	2930





Final MSE: 0.09419795221843004

Final RMSE: 0.3069168490298798

Final MAE: 0.09419795221843004

## Conclusion

After applying the GridSearch method, it can be seen that XG Boost performs slightly better than the previous default parameters as well as hypertuned Random Forest. The classification and regression performance of the model has increased by 0.01 - 0.02 in each metrics

## Neural Network Attempt (IGNORE)

```
In [ ]: '''
Artificial Neural Network Class
'''
class ANN:
    np.random.seed(10)

    '''
    Initialize the ANN;
    HiddenLayer vector : will contain the Layers' info
```

```

w, b, phi = (empty) arrays that will contain all the w, b and activat
mu = cost function
eta = a standard learning rate initialization. It can be modified by
'''
def __init__(self) :
    self.HiddenLayer = []
    self.w = []
    self.b = []
    self.phi = []
    self.mu = []
    self.eta = 1 #set up the proper Learning Rate!!

'''
add method: to add layers to the network
'''
def add(self, lay = (4, 'ReLU') ):
    self.HiddenLayer.append(lay)

'''
FeedForward method: as explained before.
'''
@staticmethod
def FeedForward(w, b, phi, x):
    return phi(np.dot(w, x) + b)

'''
BackPropagation algorithm implementing the Gradient Descent
'''
def BackPropagation(self, x, z, Y, w, b, phi):
    self.delta = []

    # We initialize auxiliar w and b that are used only inside the ba
    self.W = []
    self.B = []

    # We start computing the LAST error, the one for the OutPut Layer
    self.delta.append( (z[len(z)-1] - Y) * phi[len(z)-1](z[len(z)-1]

    '''Now we BACKpropagate'''
    # We thus compute from next-to-last to first
    for i in range(0, len(z)-1):
        self.delta.append( np.dot( self.delta[i], w[len(z)- 1 - i] ) )

    # We have the error array ordered from last to first; we flip it
    self.delta = np.flip(self.delta, 0)

    # Now we define the delta as the error divided by the number of t
    self.delta = self.delta/self.X.shape[0]

    '''GRADIENT DESCENT'''
    # We start from the first layer that is special, since it is conn
    self.W.append( w[0] - self.eta * np.kron(self.delta[0], x).reshap
    self.B.append( b[0] - self.eta * self.delta[0] )

```

```

# We now descend for all the other Hidden Layers + OutPut Layer
for i in range(1, len(z)):
    self.W.append( w[i] - self.eta * np.kron(self.delta[i], z[i-1])
    self.B.append( b[i] - self.eta * self.delta[i] )

# We return the descended parameters w, b
return np.array(self.W), np.array(self.B)

...

Fit method: it calls FeedForward and Backpropagation methods
...

def Fit(self, x_train, y_train):
    print('Start fitting...')
    ...

    Input layer
    ...

    self.X = x_train
    self.Y = y_train

    ...

    We now initialize the Network by retrieving the Hidden Layers and
    ...

    print('Model recap: \n')
    print('You are fitting an ANN with the following amount of layers

    for i in range(0, len(self.HiddenLayer)) :
        print('Layer ', i+1)
        print('Number of neurons: ', self.HiddenLayer[i][0])
        if i==0:
            # We now try to use the He et al. Initialization from ArX
            self.w.append( np.random.randn(self.HiddenLayer[i][0] , s
            self.b.append( np.random.randn(self.HiddenLayer[i][0])/np
            # Old initialization
            #self.w.append(2 * np.random.rand(self.HiddenLayer[i][0]
            #self.b.append(np.random.rand(self.HiddenLayer[i][0]))

            # Initialize the Activation function
            for act in Activation_function.list_act():
                if self.HiddenLayer[i][1] == act :
                    self.phi.append(Activation_function.get_act(act))
                    print('\tActivation: ', act)

        else :
            # We now try to use the He et al. Initialization from ArX
            self.w.append( np.random.randn(self.HiddenLayer[i][0] , s
            self.b.append( np.random.randn(self.HiddenLayer[i][0])/np
            # Old initialization
            #self.w.append(2*np.random.rand(self.HiddenLayer[i][0] ,
            #self.b.append(np.random.rand(self.HiddenLayer[i][0]))

            # Initialize the Activation function
            for act in Activation_function.list_act():
                if self.HiddenLayer[i][1] == act :

```

```

        self.phi.append(Activation_function.get_act(act))
        print('\tActivation: ', act)

    """
    Now we start the Loop over the training dataset
    """
    for I in range(0, self.X.shape[0]): # loop over the training set
        """
        Now we start the feed forward
        """
        self.z = []

        self.z.append( self.FeedForward(self.w[0], self.b[0], self.ph

        for i in range(1, len(self.HiddenLayer)): #Looping over layer
            self.z.append( self.FeedForward(self.w[i] , self.b[i], se

        """
        Here we backpropagate
        """
        self.w, self.b = self.BackPropagation(self.X[I], self.z, sel

        """
        Compute cost function
        """
        self.mu.append(
            (1/2) * np.dot(self.z[len(self.z)-1] - self.Y[I], self.z[

        )

    print('Fit done. \n')

    """
    predict method
    """
    def predict(self, x_test):

        print('Starting predictions...')

        self.pred = []
        self.XX = x_test

        for I in range(0, self.XX.shape[0]): # loop over the training set
            """
            Now we start the feed forward
            """
            self.z = []

            self.z.append(self.FeedForward(self.w[0] , self.b[0], self.ph

            for i in range(1, len(self.HiddenLayer)) : # loop over the la

```

```

        self.z.append( self.FeedForward(self.w[i] , self.b[i], se

        # Append the prediction;
        # We now need a binary classifier; we this apply an Heaviside
        # if y < 0.5 the output is zero, otherwise is zero
        self.pred.append( np.heaviside( self.z[-1] - 0.5, 1)[0] ) #

    print('Predictions done. \n')

    return np.array(self.pred)

...
We need a method to retrieve the accuracy for each training data to f
...
def get_accuracy(self):
    return np.array(self.mu)
# This is the averaged version
def get_avg_accuracy(self):
    import math
    self.batch_loss = []
    for i in range(0, 10):
        self.loss_avg = 0
        # To set the batch in 10 element/batch we use math.ceil metho
        # int(math.ceil((self.X.shape[0]-10) / 10.0)) - 1
        for m in range(0, (int(math.ceil((self.X.shape[0]-10) / 10.0)
            #self.loss_avg += self.mu[60*i+m]/60
            self.loss_avg += self.mu[(int(math.ceil((self.X.shape[0]-
        self.batch_loss.append(self.loss_avg)
    return np.array(self.batch_loss)

...
Method to set the learning rate
...
def set_learning_rate(self, et=1):
    self.eta = et

...
layers class
...
class layers :
    ...
    Layer method: used to call standar layers to add.
    Easily generalizable to more general layers (Pooling and Convolutiona
    ...
    def layer(p=4, activation = 'ReLU'):
        return (p, activation)

...
Activation functions class
...
class Activation_function(ANN):

    def __init__(self) :
```

```

    super().__init__()

    """
    Define the sigmoid activator; we ask if we want the sigmoid or its de
    """
    def sigmoid_act(x, der=False):
        if (der==True) : #derivative of the sigmoid
            f = 1/(1+ np.exp(- x))*(1-1/(1+ np.exp(- x)))
        else : # sigmoid
            f = 1/(1+ np.exp(- x))
        return f

    """
    Define the Rectifier Linear Unit (ReLU)
    """
    def reLU_act(x, der=False):
        if (der == True): # the derivative of the ReLU is the Heaviside T
            f = np.heaviside(x, 1)
        else :
            f = np.maximum(x, 0)
        return f

    def list_act():
        return ['sigmoid', 'ReLU']

    def get_act(string = 'ReLU'):
        if string == 'ReLU':
            return Activation_function.reLU_act
        elif string == 'sigmoid':
            return Activation_function.sigmoid_act
        else :
            return Activation_function.sigmoid_act

```

```

In [ ]: model = ANN()

model.add(layers.layer(8, 'ReLU'))
model.add(layers.layer(4, 'ReLU'))
model.add(layers.layer(1, 'sigmoid'))

model.set_learning_rate(0.8)

model.Fit(x_train, y_train)
acc_val = model.get_accuracy()
acc_avg_val = model.get_avg_accuracy()

predictions = model.predict(x_test)

```

Start fitting...

Model recap:

You are fitting an ANN with the following amount of layers: 3

Layer 1

Number of neurons: 8

Activation: ReLU

Layer 2

Number of neurons: 4

Activation: ReLU

Layer 3

Number of neurons: 1

Activation: sigmoid

```
-----
-
KeyError                                Traceback (most recent call last)
File ~/anaconda3/envs/capstone/lib/python3.10/site-packages/pandas/core/indexes/base.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/hashtable_class_helper.pxi:2606, in pandas._libs.hashtable.Int64HashTable.get_item()

File pandas/_libs/hashtable_class_helper.pxi:2630, in pandas._libs.hashtable.Int64HashTable.get_item()

KeyError: 0
```

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
Cell In[98], line 9
      5 model.add(layers.layer(1, 'sigmoid'))
      7 model.set_learning_rate(0.8)
----> 9 model.Fit(x_train, y_train)
     10 acc_val = model.get_accuracy()
     11 acc_avg_val = model.get_avg_accuracy()

Cell In[97], line 139, in ANN.Fit(self, x_train, y_train)
     133     self.z.append( self.FeedForward(self.w[i] , self.b[i], self.phi[i], self.z[i-1] ) )
     136 '''
     137 Here we backpropagate
     138 '''
--> 139 self.w, self.b = self.BackPropagation(self.X[I], self.z, self.Y[I])
```

```

I], self.w, self.b, self.phi)
141 '''
142 Compute cost function
143 '''
144 self.mu.append(
145     (1/2) * np.dot(self.z[len(self.z)-1] - self.Y[I], self.z[len(s
elf.z)-1] - self.Y[I])
146 )

```

File ~/anaconda3/envs/capstone/lib/python3.10/site-packages/pandas/core/series.py:1112, in Series.\_\_getitem\_\_(self, key)

```

1109     return self._values[key]
1111 elif key_is_scalar:
-> 1112     return self._get_value(key)
1114 # Convert generator to list before going through hashable part
1115 # (We will iterate through the generator there to check for slice
s)
1116 if is_iterator(key):

```

File ~/anaconda3/envs/capstone/lib/python3.10/site-packages/pandas/core/series.py:1228, in Series.get\_value(self, label, takeable)

```

1225     return self._values[label]
1227 # Similar to Index.get_value, but we do not fall back to positiona
l
-> 1228 loc = self.index.get_loc(label)
1230 if is_integer(loc):
1231     return self._values[loc]

```

File ~/anaconda3/envs/capstone/lib/python3.10/site-packages/pandas/core/indexes/base.py:3812, in Index.get\_loc(self, key)

```

3807     if isinstance(casted_key, slice) or (
3808         isinstance(casted_key, abc.Iterable)
3809         and any(isinstance(x, slice) for x in casted_key)
3810     ):
3811         raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
3813 except TypeError:
3814     # If we have a listlike key, _check_indexing_error will raise
3815     # InvalidIndexError. Otherwise we fall through and re-raise
3816     # the TypeError.
3817     self._check_indexing_error(key)

```

KeyError: 0